



Lab Recap & Docker

"What do we do with missing values?
Dockerization?"

Table of Contents

01

Kursplan &
Översikt

03

Docker

02

Pandas
Recap

04

Hosting

Overview



Moduler:

- Pandas & Missing values
- Pydantic Field Validation
- Docker
- Dockerizing
- Hosting with Supabase (bonus)
- Environment Variables recap



Utbildningsmoment

- Dataplatfformar, bakgrund och syfte
- Git och github i teamkontext ✓
- Komponenter och teknologier i en data platform ✓
- ETL vs ELT
- Utveckling av mjukvara mot databaser ✓
- Använda Python mot relationsdatabaser och andra datakällor såsom csv, http xml/json ✓
- Använda Python mot realtidsdataströmmar såsom message queues och/eller event streaming platforms
- Använda Python och för att rensa, validera och transformera data
- Workflow processer ✓

02

Pandas Recap

Transforming Data

Missing values

analytics_summary.csv

- snittpris
- medianpris
- antal produkter
- antal produkter med saknat pris

price_analysis.csv (bonus)

- top 10 dyraste produkter
- top 10 produkter med mest avvikande pris

rejected_products.csv (bonus)

- Inkludera alla produkter som är omöjliga

På ytan ser detta väldigt enkelt ut

Men det finns ett par grå zoner och
otydligheter som gör att uppgiften
blir svår...

VI UTFORSKAR

Impossible Values (Avvisa (REJECT))



Är något 'omöjligt' att bearbeta utan att förändra värdet (risk för felaktig data) så avvisas det.


Ex:

- negativa priser (tar vi bort - så blir det helt plötsligt ett positivt värde, matchar det? Omöjligt att ta reda på)
- `Created_at` datum som ligger efter `updated_at`. Omöjligt att ta reda på
- ID som är borttagna. Hur bearbetas detta? Finns det koppling till Primary_Key i databas? Om inte, omöjligt att ta reda på.

Medan överdrivna priser är prima exempel på data som oftast bör flaggas.

- Gold_plates_shoes: 7200kr (makes sense)
- bananas: 4828kr (weird - FLAG)

When do we REJECT? (Information == Värde)



Hur ska vi egentligen avgöra när data är "omöjlig" och ska avvisas, och när den bara är osäker och bör flaggas?

Till exempel: om `currency` är NaN – är det automatiskt en reject, eller kan det finnas scenarier där vi istället flaggar?

När vi flaggar något, är syftet enbart att uppmärksamma osäkerhet för manuell granskning – eller förväntas vi också korrigera värdet om det är uppenbart (t.ex. anta att det ska vara "SEK")?

Var går gränsen mellan datavalidering och datakorrigering?

What if i'm wrong? (Felaktig data & risker)

Ett annat exempel gällande felaktig data

Organisation har ett uppdrag till oss, de frågar om en kan analysera top 10 månads intäckter -> vi glömmer att städa data -> och median fallerar för att vi råkat ha felaktig data med -> dåligt resultat -> reflekteras dåligt på oss som utvecklare



What if i'm wrong? #2

(Felaktig data & risker)

Precis!

För nu handlar det delvist om konsekvenstänkande.

Låt oss säga att vi släpper in

Sek 500

currency NONE

Om vi inte kan få reda på informationen så är det **direkt avvisande**.

Är det USD? EUR? SEK? Detta kan bli väldigt felaktig data i längden om vi stöter på fler fall. Denna grå zon är perfekt exempel på när det är viktigt att analysera organisationens behov.

I vårt fall - **avvisa direkt**

Datakvalitet kan bli affärsrisk.

Vi får aldrig gissa affärskritisk information i ett datalager.

Antaganden hör hemma i kravspecifikation (i detta fall har vi inget, därav blir det svårt, vilket är avsiktligt för denna labboration)

Här kommer ett exempel:

--PRICE LIST--

price

100
200
free
300



Om vi sedan kalkylerar .mean() på 'none'

Python

```
prices = pd.Series([100, 200, None, 300])  
prices.mean()
```



resultat på medelvärdet blir **200**. $((100 + 200 + 300) / 3 = 200)$

Varför är detta farligt?

Solution

(Felaktig data & risker)



Precis!

Svaret blir ett annat resultat: $(100 + 200 + 0 + 300) / 4 = 150$

Om data inte går att använda

Är datan omöjlig att bearbeta, omöjlig att gissa sig fram till = avvisa


None & NaN

Betyder att värdet är okänt. Om resterande av raden är giltig och innehåller värdefull information = flagga

Free

Betyder att värdet är känt, men inte numerisk representation

Data modelling? (Felaktig data & risker)



I ett verkligt scenario (*överkurs just nu*)

Innan du ens tittar på `NaN` måste du veta:

- Är pris obligatoriskt?
- Får pris vara 0?
- Får pris vara negativt?
- Får pris saknas?

Detta är ett **datakontrakt relaterad fråga**, det är därför det blir förvirrande då vi arbetar inom pandas för att försöka lösa problemet.

Det är här **datamodellering** kommer in

What if i'm wrong? #3 (Felaktig data & risker)

Ett sista exempel - Datamodellering

Datamodellering löser oftast problemet med osäkerhet om vad som får komma in genom att definiera om datan är obligatoriskt tillåten eller meningsfull

```
price_amount DECIMAL NOT NULL  
currency_code CHAR(3) NOT NULL
```

KPI vs Quality Assurance

(The Difference)

Key Performance Indicator

Mäter hur bra verksamheten presterar

- Genomsnittligt ordervärde
- Antal aktiva användare
- Andel godkända transaktioner
- Omsättning per månad

Quality Assurance

Säkerställer att datan är korrekt, komplett och tillförlitlig

- Saknade värden (NaN)
- Negativa priser
- Felaktigt datumformat
- Dubbletter

*Kvalitetssäkring är **förutsättningen** för att KPI:er ska vara trovärdiga.*

Summarized **(Information == Värde)**

Finansiell rapportering / KPI:er
oftast endast *valid data*

Explorativ analys / kvalitetssäkring
valid + flaggad (med transparens)

Pydantic Modelling



03

Docker

Docker Install & Version



Is docker installed? (command)



```
$ docker --version
```

Installation (MAC & Windows)

MAC

<https://docs.docker.com/desktop/setup/install/mac-install/>

Windows

<https://docs.docker.com/desktop/setup/install/windows-install/>

Docker

What is it?




Docker

(What is it)

Docker

An **open source platform** that allows developers and systems administrators to **package applications** into **containers**



Docker

(What is it)

What is Docker?



Copy as Markdown



Docker is an open platform for developing, shipping, and running applications. Docker enables you to separate your applications from your infrastructure so you can deliver software quickly. With Docker, you can manage your infrastructure in the same ways you manage your applications. By taking advantage of Docker's methodologies for shipping, testing, and deploying code, you can significantly reduce the delay between writing code and running it in production.

Read more: <https://docs.docker.com/get-started/docker-overview/>

Docker Terminology

Container

The Problem

Explanation

Imagine you're developing a killer web app that has three main components - a React frontend, a Python API, and a PostgreSQL database. If you wanted to work on this project, you'd have to install Node, Python, and PostgreSQL.

How do you make sure you have the same versions as the other developers on your team? Or your CI/CD system? Or what's used in production?

How do you ensure the version of Python (or Node or the database) your app needs isn't affected by what's already on your machine? How do you manage potential conflicts?

Source: <https://docs.docker.com/get-started/docker-concepts/the-basics/what-is-a-container>

What is a container? Simply put, containers are isolated processes for each of your app's components. Each component - the frontend React app, the Python API engine, and the database - runs in its own isolated environment, completely isolated from everything else on your machine.

Here's what makes them awesome. Containers are:

- Self-contained. Each container has everything it needs to function with no reliance on any pre-installed dependencies on the host machine.
- Isolated. Since containers run in isolation, they have minimal influence on the host and other containers, increasing the security of your applications.
- Independent. Each container is independently managed. Deleting one container won't affect any others.
- Portable. Containers can run anywhere! The container that runs on your development machine will work the same way in a data center or anywhere in the cloud!

Docker Terminology

Image

Image

Explanation

Seeing as a **container** is an isolated process, where does it get its files and configuration? How do you share those environments?

That's where container images come in. A container image is a standardized package that includes all of the files, binaries, libraries, and configurations to run a container.

For a PostgreSQL image, that image will package the database binaries, config files, and other dependencies. For a Python web app, it'll include the Python runtime, your app code, and all of its dependencies.

Source: <https://docs.docker.com/get-started/docker-concepts/the-basics/what-is-an-image/>

Docker Terminology

Dockerfile

Dockerfile

Dockerfile reference




Copy as Markdown



Docker can build images automatically by reading the instructions from a Dockerfile. A Dockerfile is a text document that contains all the commands a user could call on the command line to assemble an image. This page describes the commands you can use in a Dockerfile.

Source: <https://docs.docker.com/reference/dockerfile/>

Dockerfile Example



Den MÅSTE döpas till
'Dockerfile'

Dockerfile Example #2

```
FROM python:3.12-slim

WORKDIR /app

# Install uv
RUN pip install uv

# Copy only dependency files first (better layer caching)
COPY pyproject.toml uv.lock ./

# Install dependencies from lock file
RUN uv sync --frozen

# Copy rest of the project
COPY . .

EXPOSE 8000

CMD ["uv", "run", "uvicorn", "app.main:app", "--host", "0.0.0.0", "--port", "8000"]
```

Dockerfile Instructions

ONBUILD

Specify instructions for when the image is used in a build.

RUN

Execute build commands.

SHELL

Set the default shell of an image.

STOP SIGNAL

Specify the system call signal for exiting a container.

USER

Set user and group ID.

VOLUME

Create volume mounts.

WORKDIR

Change working directory.

Instruction	Description
ADD	Add local or remote files and directories.
ARG	Use build-time variables.
CMD	Specify default commands.
COPY	Copy files and directories.
ENTRYPOINT	Specify default executable.
ENV	Set environment variables.
EXPOSE	Describe which ports your application is listening on.
FROM	Create a new build stage from a base image.
HEALTHCHECK	Check a container's health on startup.
LABEL	Add metadata to an image.
MAINTAINER	Specify the author of an image.

Docker Commands

Common

Docker Commands

(List running containers)



```
$ docker ps
```



Lists all docker containers.


Add -a at the end to list stopped ones.

Docker Commands

(Generate an Image)



```
$ docker build -t demo-9 .
```



-t == tag name

. == from where we currently stand (path)

Docker Commands

(Generate an Image)



```
$ docker run -p 8000:8000 demo-9
```



Create a container with an Image inside

Expose port 8000 to the local machine

Demo-9 is the image in question, without it this command won't work!

Rekommendera att testa utan -p 8000:8000 för debugging

Docker Commands BONUS

(Generate an Image)



```
$ docker run --rm -p 8000:8000  
demo-9
```



Docker Hygiene by removing old container

Docker Visual Example With Commands




Build Image

(`docker build -t NAME .`)

This means success!

```
=> => extracting sha256:09fa645f2e8f01bf76e8432e1b0c11b79a8726f64b60501af160c71b7c3917ab
=> [internal] load build context
=> => transferring context: 53.13MB
=> [2/6] WORKDIR /app
=> [3/6] RUN pip install uv
=> [4/6] COPY pyproject.toml uv.lock .
=> [5/6] RUN uv sync --frozen
=> [6/6] COPY . .
=> exporting to image
=> => exporting layers
=> => writing image sha256:a215570c97a1ea19f258c9b3fa13eb2392f960dd3fc38c1286e01d259a175778
=> => naming to docker.io/library/demo-9
```

View build details: docker-desktop://dashboard/build/desktop-linux/desktop-linux/u7xr3e0ivmx8rwy4yt4aiy5kd



Local My Hub

0 Bytes / 765.74 MB in use 3 images Last refresh: 45 minutes ago

	Name	Tag	Image ID	Created	Size	Actions
<input type="checkbox"/>	postgres	latest	21a7e541217c	5 months ago	478.92 MB	
<input type="checkbox"/>	testcontainers/ryuk	0.12.0	3b9e208a7ab3	9 months ago	15.84 MB	
<input type="checkbox"/>	demo-9	latest	a215570c97a1	2 minutes ago	350.27 MB	

Run App

(`docker run demo-9`)

	Name	Tag	Image ID	Created	Size	Actions
<input type="checkbox"/>	<input checked="" type="radio"/> postgres	latest	21a7e541217c	5 months ago	478.92 MB	▶ ⋮ trash
<input type="checkbox"/>	<input checked="" type="radio"/> testcontainers/ryuk	0.12.0	3b9e208a7ab3	9 months ago	15.84 MB	▶ ⋮ trash
<input type="checkbox"/>	<input checked="" type="radio"/> demo-9	latest	a215570c97a1	15 minutes ag	350.27 MB	▶ ⋮ trash

Notice it's green!

Container up and running? (YES)

The screenshot shows the Docker desktop application interface. The left sidebar has a dark theme with white icons and text. The 'Containers' option is selected, indicated by a blue background. Other options include 'Ask Gordon' (BETA), 'Images', 'Volumes', 'Kubernetes', 'Builds', 'Models', 'MCP Toolkit' (BETA), 'Docker Hub', 'Docker Scout', and 'Extensions'. The main area is titled 'Containers' with a 'Give feedback' link. It displays system usage statistics: 'Container CPU usage' at 0.29% / 800% (8 CPUs available) and 'Container memory usage' at 246.6MB / 7.47GB. A search bar and a filter button 'Only show running containers' are present. A table lists the running container 'ecstatic_meitne' with details: Container ID fa078b609faa, Image demo-9, Port(s) 8000:8000, CPU (%) 0.29%, and Last status 44 seconds ago. Action buttons for the container row are shown.

	Name	Container ID	Image	Port(s)	CPU (%)	Last status	Actions
<input type="checkbox"/>	ecstatic_meitne	fa078b609faa	demo-9	8000:8000	0.29%	44 sec	Running ⋮ Delete

Notice it has exposed port 8000!



nice_proskuriakova

📦 d987f3c30b69 ⚡ ⚡ demo-9:latest

STATUS

Running (4 minutes ago)



Logs

Inspect

Bind mounts

Exec

Files

Stats

```
warning: Ignoring existing virtual environment linked to non-existent Python interpreter: .venv/bin/python3 -> python
Downloading cpython-3.9.25-linux-aarch64-gnu (download) (27.6MiB)
Downloaded cpython-3.9.25-linux-aarch64-gnu (download)
Using CPython 3.9.25
Removed virtual environment at: .venv
Creating virtual environment at: .venv
Downloading psycopg-binary (4.4MiB)
Downloading uvloop (3.5MiB)
Downloading pydantic-core (1.8MiB)
Downloaded pydantic-core
Downloaded uvloop
Downloaded psycopg-binary
Installed 47 packages in 373ms
INFO:     Started server process [48]
INFO:     Waiting for application startup.
INFO:     Application startup complete.
INFO:     Uvicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)
```



So... What's the big deal? Explained



Container up and running? (YES)

Reproducible environment

Same Python, same OS, same dependencies... everywhere!

No “works on my machine” problems

What runs locally runs the same in production.

Isolation

No dependency conflicts, no global pollution, easy cleanup.

Deployable unit

The app + runtime + config are packaged into one portable artifact.


04

Hosting (bonus)
& Docker Env variables

.env
FIX

```
curl_setopt(conn, CURLOPT_URL, url);
if (curl_error(conn) != CURLE_OK)
    error("Failed to set URL [%s]\n", curl_error(conn));
else
    curl_easy_setopt(conn, CURLOPT_FOLLOWLOCATION,
                     1);
if (curl_error(conn) != CURLE_OK)
    error("Failed to set redirect option [%s]\n",
          curl_error(conn));
else
    curl_easy_setopt(conn, CURLOPT_WRITEFUNCTION,
                     writeFunction);
if (curl_error(conn) != CURLE_OK)
    error("Failed to set writer [%s]\n",
          curl_error(conn));
```

Hosting PostgreSQL



Build in a weekend Scale to millions

Supabase is the Postgres development platform.

Start your project with a Postgres database, Authentication, instant APIs, Edge Functions, Realtime subscriptions, Storage, and Vector embeddings.

\$0 / month

Supabase (Free Tier)

Get started with:

- ✓ Unlimited API requests
- ✓ 50,000 monthly active users
- ✓ 500MB database size
Shared CPU • 500MB RAM
- ✓ 5 GB egress
- ✓ 5 GB cached egress
- ✓ 1 GB file storage
- ✓ Community support

Read more: <https://supabase.com/pricing>

GitHub + Supabase (Connect to Account)

<https://supabase.com/>

Sign up using GitHub for easy repository access

Create Organization (HUB for projects)

Your Organizations

Search for an organization

+ New organization



krillinator
Free Plan



Krillinator's Org
Free Plan

Create a new project

Your project will have its own dedicated instance and full Postgres database.

An API will be set up so you can easily interact with your new database.

Organization

krillinator FREE

Project name

demo-9

Database password

••••••••••••••

 Copy

This password is strong. Generate a password.

Region

 Europe

Select the region closest to your users for the best performance.

Security

Enable Data API

Autogenerate a RESTful API for your public schema.
Recommended if using a client library like [supabase-js](#).



Client libraries need Data API to query your database

Disabling it means supabase-js and similar libraries
can't query or mutate data.

ADVANCED CONFIGURATION >

Cancel

Create new project

New Project #2 (Disable Data API)

We disable it because we don't want to use it against a frontend UI.

We're using **psycopg3** with Python instead



krillinator

FREE

Project Overview

Table Editor

SQL Editor

Database

Authentication

Storage

Edge Functions

Realtime

Advisors

Observability

Logs

Integrations

Project Settings

Create First Table (Supabase)

The screenshot shows the Supabase Table Editor interface. On the left, there's a sidebar with various project management and database-related tabs. The main area is titled "Table Editor" and shows a list of tables under the schema "public". A prominent red arrow points from the top right towards a call-to-action button labeled "Create a table".

Table Editor

schema public

+ New table

Search tables...

No tables or views
Any tables or views you create will be listed here.

- postgres_table_0
- postgres_table_1
- postgres_table_2
- postgres_table_3

Create a table
Design and create a new database table

Recent items

No recent items yet
Items will appear here as you browse through your project

Create First Table #2

(Payload: jsonB)

Columns

About data types

Import data from CSV

Name ?

Type

Default Value ?

Primary

≡

id



int8



NULL



1



≡

created_at



timestamp



now()



≡

payload



{ } jsonb



NULL



Notice: You can also import data from CSV!

[Project Overview](#)[Table Editor](#)[SQL Editor](#)[Database](#)[Authentication](#)[Storage](#)[Edge Functions](#)[Realtime](#)[Advisors](#)[Observability](#)[Logs](#)[Integrations](#)[Project Settings](#)


schema public

ENT


BONUS: ERD table

raw_events	
♂	# id int8
◆	created_at timestampz
◆	payload jsonb

Connection String IPv4



Create First Table (Supabase)



Connect to your project

Get the connection strings and environment variables for your app.

Connection String App Frameworks Mobile Frameworks ORMs API Keys MCP

Session pooler works for IPv4

Type

URI 

Source

Primary Database 

Method

Session pooler 

 Learn how to connect to your Postgres databases. [Read docs](#) 

Session pooler  SHARED POOLER

Only recommended as an alternative to Direct Connection, when connecting via an IPv4 network.

postgresql://postgres.pfjscancbcuhplqbvyos:[YOUR-PASSWORD]@aws-1-eu-west-1.pooler.supabase.com:5432/postgres

 View parameters



IPv4 compatible

Session pooler connections are IPv4 proxied for free



Only use on a IPv4 network

Session pooler connections are IPv4 proxied for free.
Use Direct Connection if connecting via an IPv6 network.

Copy

Environment Variables Step by Step



Environment Variables (Dependency)

```
$ pip install python-dotenv  
$ uv add python-dotenv # for uv
```


Environment Variables

(File Creation)



```
$ touch .env
```

Environment Variables (Results)



.gitignore (example)

```
∅ .gitignore          9      # Virtual environments
≡ .python-version    10     .venv
[T] pyproject.toml   11     .env
M README.md
[T] uv.lock
```

.env (Content)



A screenshot of a terminal window with a dark background. At the top, there are three tabs: 'main.py' (highlighted in yellow), '.gitignore' (greyed out), and '.env' (highlighted in blue). The '.env' tab has a close button ('x'). Below the tabs, the content of the '.env' file is displayed in a monospaced font:

```
1 DB_PORT="5432"
2 DB_PASSWORD="Benny_Banana_4"
3 DB_USERNAME="postgres"
4 DB_HOST="aws-1-eu-west-1.pooler.supabase.com"
5 DB_NAME="pfjscancbcuhplqbvyos"
```

Environment Variables (Results)

```
from dotenv import load_dotenv
import os

load_dotenv()    # reads .env from project root

DB_PORT = os.getenv("DB_PORT")
DB_USERNAME = os.getenv("DB_USERNAME")
DB_PASSWORD = os.getenv("DB_PASSWORD")
DB_NAME = os.getenv("DB_NAME")
DB_HOST = os.getenv("DB_HOST")
DATABASE_URL =
f"postgresql://{{DB_USERNAME}}.{{DB_NAME}}:{{DB_PASSWORD}}@{{DB_HOST}}:{{DB_PORT}}/postgres"
```

Function Example (Insert)

```
@app.post("/product", status_code=status.HTTP_201_CREATED, response_model=Product)
def post_product(product: Product) -> Product:

    with pool.connection() as conn:
        with conn.transaction():
            conn.execute(
                "INSERT INTO raw_events (payload) VALUES (%s)" ,
                (json.dumps(product.model_dump()),)
            )

    return product
```

Results (The End)

Table Editor

raw_events

schema public

+ New table

Search tables...

raw_events

	<code>id</code> int8	<code>created_at</code> timestampz	<code>payload</code> jsonb	
1		2026-02-22 19:50:54.618905+0	{"name": "Bananas", "price": 150, "quantity": 15}	
2		2026-02-22 19:52:28.675482+0	{"name": "Bananas", "price": 150, "quantity": 15}	
3		2026-02-22 19:53:53.207811+0	{"name": "apples", "price": 190, "quantity": 13}	

Docker

.env

.dockerIgnore

(gitignore but: for Docker)

1		.env
2	📁	.venv
3		__pycache__
4	📁	.git
5		

A screenshot of a file browser interface, likely PyCharm, showing a directory structure. The left pane displays a tree view with a expanded folder named 'demo-9'. Inside 'demo-9', there are two subfolders: '.venv' and 'app'. Below 'app' is a file named '.dockerignore', which is highlighted with a dark gray rectangle and features a white circle with a diagonal slash icon over it. The right pane shows a list of five items, each consisting of a number (1 through 5), a file icon, and a file name: '.env', '📁 .venv', '__pycache__', '📁 .git', and an empty row (5). The background of the right pane has a subtle grid pattern.

Is docker installed? (command)



```
$ docker run --env-file .env -p  
8000:8000 fastapi-app
```

Injicera .env filen!

Is docker installed? (command)

```
$ docker run --rm -it fastapi-app  
sh -lc "ls -la /app | grep .env ||  
echo 'No secrets exposed: OK'"
```

Verifiera att du inte råkat skriva in dina hemligheter

THANKS !

Do you have any questions?
kristoffer.johansson@sti.se

CREDITS: This presentation template was created by Slidesgo, including icons by Flaticon, and infographics & images by Freepik.