



Docker Compose & Kafka

"Crafting a functional pipeline for realtime data streaming"

Table of Contents

01

Kursplan &
Översikt

03

Kafka

02

Docker Compose

Overview



Moduler:

- Docker Compose
- Kafka (*and terminologies*)
- Kafka Consumers (*Bonus*)



Utbildningsmoment

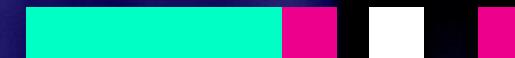
- Dataplatfformar, bakgrund och syfte
- Git och github i teamkontext ✓
- Komponenter och teknologier i en data platform ✓
- ETL vs ELT
- Utveckling av mjukvara mot databaser ✓
- Använda Python mot relationsdatabaser och andra datakällor såsom csv, http xml/json ✓
- Använda Python mot realtidsdataströmmar såsom message queues och/eller event streaming platforms ✓
- Använda Python och för att rensa, validera och transformera data
- Workflow processer ✓

02

Docker Compose

Multi-containers

Why?



Docker Compose is a tool for defining and running multi-container applications. It is the key to unlocking a streamlined and efficient development and deployment experience.

Compose simplifies the control of your entire application stack, making it easy to manage services, networks, and volumes in a single YAML configuration file. Then, with a single command, you create and start all the services from your configuration file.

Compose works in all environments - production, staging, development, testing, as well as CI workflows. It also has commands for managing the whole lifecycle of your application:


- Start, stop, and rebuild services
- View the status of running services
- Stream the log output of running services
- Run a one-off command on a service

Docker Compose (Multi-containers)

När används det?

- Development Environments (*databases, queues, caches, web service APIs, etc*)
- Automated Testing Environments (*creates & destroys containers*)
- Multiple services, one host


Läs mer: <https://docs.docker.com/compose/>





Private Network

Attaches all service containers to the same private network



```
DATABASE_URL = "postgresql://postgres:password@db:5432/appdb"
```

@db = hostname (*DNS*)

```
services:  
  api:  
    build: .  
  db:  
    image: postgres:16
```

A service is an abstract definition of a computing resource within an application which can be scaled or replaced independently from other components. Services are backed by a set of containers, run by the platform according to replication requirements and placement constraints. As services are backed by containers, they are defined by a Docker image and set of runtime arguments. All containers within a service are identically created with these arguments.

A Compose file must declare a `services` top-level element as a map whose keys are string representations of service names, and whose values are service definitions. A service definition contains the configuration that is applied to each service container.

Läs mer <https://docs.docker.com/reference/compose-file/services/>

Example of what a **collection** of **services** look like:

```
version: "3.9"

services:
  api:
    build: .
    ports:
      - "8000:8000"
    depends_on:
      - db
    environment:
      DATABASE_URL: postgres://postgres:postgres@db:5432/appdb

  db:
    image: postgres:16
    environment:
      POSTGRES_USER: postgres
      POSTGRES_PASSWORD: postgres
      POSTGRES_DB: appdb
```

Download Project Git Clone



Clone the repository

Open terminal

```
git clone <repo-url>
cd <project-folder>
```

Install dependencies

From the project root:

```
uv sync
```

This creates a virtual environment and installs all dependencies from `uv.lock`.

Create a file named `.env` in the project root : Or checkout the `.env-example` file for quicker setup

```
DB_USER=postgres
DB_PASSWORD=postgres
DB_HOST=localhost
DB_PORT=5432
DB_NAME=products
```


<https://github.com/Krillinator/data-platform-lektion-10-kafka-docker-compose/tree/main>

Docker Compose Step by Step



Docker Compose (FastAPI + PostgreSQL)


 FastAPI



NOTERA: Vi behöver inte PostgreSQL installerat eller öppet här!

För att: Den kommer skapas av en 'image'. Dockers superkraft!

Docker Compose (File)



Kan också heta 'docker-compose'



```
services:  
  
postgres: # must match .env hostname  
  image: postgres:18  
  container name: postgres  
  environment:  
    POSTGRES_USER: ${DB_USER}  
    POSTGRES_PASSWORD: ${DB_PASSWORD}  
    POSTGRES_DB: ${DB_NAME}  
  volumes:  
    - postgres_data:/var/lib/postgresql  
  
app:  
  build: .  
  env file:  
    - .env  
  depends on:  
    - postgres  
  ports:  
    - "8000:8000"  
  
volumes:  
  postgres_data:
```

Docker Compose (Content)

Docker containers are 'ephemeral'
vilket innebär att data tas bort.

'Volumes' vidbehåller data

Docker Compose (Build)




```
$ docker compose up --build
```

This will read the Dockerfile, build the image and start the containers.

```
(.venv) kristofferjohansson@Kristoffers-MacBook-Pro lektion_10 % docker compose up --build
[+] up 12/14
.. Image postgres:18 [██████████] 109MB / 160.9MB Pulling
```

Docker Compose (up and running)



The screenshot shows the Docker Desktop interface. At the top, there are three colored dots (red, orange, green) followed by the "docker desktop" logo and the word "PERSONAL". Below this is a table with the following data:

	Name	Tag	Image ID	Created	Size	Actions
<input type="checkbox"/>	postgres	latest	21a7e541217c	5 months ago	478.92 MB	▶ ⋮ trash
<input type="checkbox"/>	testcontainers/ryuk	0.12.0	3b9e208a7ab3	9 months ago	15.84 MB	▶ ⋮ trash
<input checked="" type="checkbox"/>	postgres	18	50a72b3e27d0	23 hours ago	479.06 MB	▶ ⋮ trash
<input type="checkbox"/>	<none>	<none>	dc8c66259070	3 minutes ago	298.92 MB	▶ ⋮ trash
<input checked="" type="checkbox"/>	lektion_10-app	latest	db400cff71ea	1 second ago	298.92 MB	▶ ⋮ trash

- Ask Gordon BETA
- Containers ←
- Images
- Volumes
- Kubernetes
- Builds
- Models
- MCP Toolkit BETA
- Docker Hub
- Docker Scout

Extensions



lektion_10

/Users/kristofferjohansson/PycharmProjects/lektion_10

[View configurations](#)



postgres •
postgres:18



app •
lektion_10-app
8000:8000 ↗

app



down at 2026-02-25 18:13:33 UTC

2026-02-25 18:13:33.421 UTC [1] LOG: database system is ready to accept connections

Downloaded cpython-3.9.25-linux-aarch64-gnu (download)

Using CPython 3.9.25

Removed virtual environment at: .venv

Creating virtual environment at: .venv

Downloading psycopg-binary (4.4MiB)

Downloading pydantic-core (1.8MiB)

Downloading uvloop (3.5MiB)

Downloaded pydantic-core

Downloaded uvloop

Downloaded psycopg-binary

Installed 47 packages in 397ms

INFO: Started server process [54]

INFO: Waiting for application startup.

INFO: Application startup complete.

INFO: Uvicorn running on <http://0.0.0.0:8000> ↗ (Press CTRL+C to quit)

INFO: 192.168.65.1:47542 - "GET / HTTP/1.1" 404 Not Found

INFO: 192.168.65.1:31837 - "GET /favicon.ico HTTP/1.1" 404 Not Found

INFO: 192.168.65.1:47542 - "GET / HTTP/1.1" 404 Not Found

INFO: 192.168.65.1:31837 - "GET /favicon.ico HTTP/1.1" 404 Not Found



POST



http://localhost:8000/products

Params

Auth

Headers (8)

Body

Pre-req.

Tests

Settings

raw

JSON

```
1  {
2    "name": "Bananas",
3    "price": 22,
4    "quantity": 5
5 }
```

GET



localhost:8000/products

Params Auth Headers (6) Body Pre-req. Tests Settings

Body



200 OK

4

Pretty

Raw

Preview

Visualize

JSON



```
1  []
2  {
3      "name": "Bananas",
4      "price": 22.0,
5      "quantity": 5,
6      "id": 1,
7      "created_at": "2026-02-25T18:25:47.293518Z"
8  },
9  {
10     "name": "Bananas",
11     "price": 22.0,
12     "quantity": 5,
13     "id": 2,
14     "created_at": "2026-02-25T18:28:53.430523Z"
15 }
16 []
```

Detta kommer ligga kvar även under omstart

03


Kafka

Kafka

What is it?

Apache Kafka

Apache Kafka is an open-source distributed event streaming platform used by thousands of companies for high-performance data pipelines, streaming analytics, data integration, and mission-critical applications.



Above is a snapshot of the number of top-ten largest companies using Kafka, per-industry.

Läs mer: <https://kafka.apache.org/>

Core capabilities

Kafka boasts core capabilities that are battle tested and ready to power businesses in the digital world.



High Throughput

Deliver messages at network limited throughput using a cluster of machines with latencies as low as 2ms.



Scalable

Scale production clusters up to a thousand brokers, trillions of messages per day, petabytes of data, hundreds of thousands of partitions. Elastically expand and contract storage and processing.



Permanent Storage

Store streams of data safely in a distributed, durable, fault-tolerant cluster.



High Availability

Stretch clusters efficiently over availability zones or connect separate clusters across geographic regions.



Built-in Stream Processing

Process streams of events with joins, aggregations, filters, transformations, and more, using event-time and exactly-once processing.



Connect to almost anything


Kafka's out-of-the-box Connect interface integrates with hundreds of event sources and event sinks including Postgres, JMS, Elasticsearch, AWS S3, and more.

History (bonus) (Zookeeper vs KRaft)

ZooKeeper is a **standalone distributed system** used for coordination.

Separate installation that was ran as its own cluster.

Other systems (like Kafka) connected to it. (*Detta innebär att det inte var integrerat*)



History (bonus) #2

(Zookeeper vs KRaft)

It's a **coordination service**.

It stores small pieces of shared metadata like:

- Who is leader?
- Which nodes are alive?
- What configuration exists?
- What topics exist (in Kafka's case)?




Apache ZooKeeper™




Apache ZooKeeper™

ZooKeeper: Ett separat distribuerat koordineringssystem

Zookeeper



KRaft



Läs mer <https://kafka.apache.org/42/operations/kraft/>

Kafka & Coordination

(Why is it needed?)

För att Kafka är ett **distribuerat system**.

- Distribuerat = flera servrar samtidigt.

Så fort du har flera servrar uppstår problem

Kafka Terminologies



Kafka Terms

(Record)

En **record** är: Ett enskilt meddelande

Exempelvis:

```
</> JSON  
  
{ "order_id": 123, "amount": 500 }
```



Kafka Terms

(Records are sent to Topics)

Record



```
{ order_id: 5, price: 50 }
```

Record



```
{ order_id: 3, price: 23 }
```



Kafka Terms (Partition)

Topic: orders

- └ Partition 0
- └ Partition 1
- └ Partition 2

Partition är en strikt 'append only'
log

En partition innehåller flera Records

Partition 0

Offset 0 → {order A}

Offset 1 → {order B}

Offset 2 → {order C}

Topic: orders

Partition 0

offset 0 → {order 123}

offset 1 → {order 124}

Partition 1

offset 0 → {order 200}

offset 1 → {order 201}

Sammanfatta (Termer)

Broker & Controller (bonus)

What is it?

Kafka Broker

(What is it?)

A server process that stores topic data, handles client requests (producers/consumers), and replicates partitions.

Kafka brokers:

- Accepts messages from producers (*producers is a client app that sends a message (topic)*)
- Stores them on disk
- Serves messages to consumers
- Replicates data to other brokers
- Hosts topic partitions
- May be leader or follower for partitions

Kafka Controller (What is it?)

Controllern styr **metadata**, inte data. Den bestämmer:

- Vem ska vara leader för partition 0?
- Vad händer om Broker 1 dör?
- Hur ska partitions fördelas?
- När skapas topics?

Den hanterar **klusterlogik**, inte records.

Configuration

Process Roles

In KRaft mode each Kafka server can be configured as a controller, a broker, or both using the `process.roles` property. This property can have the following values:

- If `process.roles` is set to `broker`, the server acts as a broker.
- If `process.roles` is set to `controller`, the server acts as a controller.
- If `process.roles` is set to `broker,controller`, the server acts as both a broker and a controller.

Kafka servers that act as both brokers and controllers are referred to as “combined” servers. Combined servers are simpler to operate for small use cases like a development environment. The key disadvantage is that the controller will be less isolated from the rest of the system. For example, it is not possible to roll or scale the controllers separately from the brokers in combined mode. Combined mode is not recommended in critical deployment environments.

Kafka Docker Compose



Explore / bitnami / kafka



bitnami/kafka



Verified Publisher

By [VMware](#) · Updated 6 months ago

Bitnami Secure Image for kafka

INTEGRATION & DELIVERY

MESSAGE QUEUES

MONITORING & OBSERVABILITY

☆ 959 ↓ 100M+

Image: <https://hub.docker.com/r/bitnami/kafka>

compose.yml

```
kafka:  
  image: apache/kafka:4.1.1  
  container_name: kafka  
  ports:  
    # host -> kafka (for tools running on your laptop)  
    - "29092:29092"  
  volumes:  
    - kafka data:/bitnami  
  environment:  
    KAFKA_NODE_ID: "1"  
    KAFKA_PROCESS_ROLES: "broker,controller"  
    KAFKA_CONTROLLER_QUORUM_VOTERS: "1@kafka:9093"  
    KAFKA_LISTENERS: "PLAINTEXT://:9092,CONTROLLER://:9093"  
    KAFKA_ADVERTISED_LISTENERS: "PLAINTEXT://kafka:9092"  
    KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: "PLAINTEXT:PLAINTEXT,CONTROLLER:PLAINTEXT"  
    KAFKA_CONTROLLER_LISTENER_NAMES: "CONTROLLER"  
    KAFKA_INTER_BROKER_LISTENER_NAME: "PLAINTEXT"  
    KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: "1"  
    KAFKA_TRANSACTION_STATE_LOG_REPLICATION_FACTOR: "1"  
    KAFKA_TRANSACTION_STATE_LOG_MIN_ISR: "1"
```


compose.yml

```
volumes:  
  postgres_data:  
  kafka_data:
```



Uppdatera 'volumes' & depends 'on'

```
app:  
  build: .  
  env_file:  
    - .env  
  depends_on:  
    postgres:  
      condition: service_started  
    kafka:  
      condition: service_started  
  
  ports:  
    - "8000:8000"
```



Kafka UI Bonus




The screenshot shows the UI for Apache Kafka version 0.7.2 at http://localhost:8080. The left sidebar has a 'local' node selected. The main dashboard shows one online cluster ('local') with 1 broker, 0 partitions, and 0 topics, all producing 0 bytes. There is a toggle switch for 'Only offline clusters' which is off.

Cluster name	Version	Brokers count	Partitions	Topics	Production	Consumption
local	1.0-UNKNOWN	1	0	0	0 Bytes	0 Bytes

Bra för debugging och visuell feedback

Kafka cluster = en/flera Kafka brokers som arbetar tillsammans

Kafka Streams With Products



Kafka Python (synchronous) (install)



```
$ uv add kafka-python
```

Injicera .env filen!

.env (Adding topics & servers)

≡ .env

```
KAFKA_BOOTSTRAP_SERVERS=kafka:9092
PRODUCTS_TOPIC=products.created
```

Main.py

(Loading Kafka .env variables)

```
# Kafka Setup
KAFKA_BOOTSTRAP_SERVERS = os.getenv("KAFKA_BOOTSTRAP_SERVERS",
"kafka:9092")
PRODUCTS_TOPIC = os.getenv("PRODUCTS_TOPIC", "products.created")
```

Main.py (Setup Producer)

 main.py

```
# Kafka producer (sync)
app.state.kafka_producer = KafkaProducer (
    bootstrap_servers=KAFKA_BOOTSTRAP_SERVERS,
    value_serializer=lambda v: json.dumps(v).encode("utf-8"),
    key_serializer=lambda k: k.encode("utf-8") if k else None,
)

yield

# Shutdown
try:
    app.state.kafka_producer.close()
except Exception:
    pass

app.state.pool.close()
```

Main.py (Setup Producer)

 main.py

```
# Publish "product.created" event to Kafka
event = {
    "type": "product.created",
    "product_id": row["id"],
    "name": row["name"],
    "price": str(row["price"]),
    "quantity": row["quantity"],
    "created_at": row["created_at"].isoformat(),
}

app.state.kafka_producer.send (
    PRODUCTS_TOPIC,
    key=str(row["id"]),
    value=event,
)
```

POST



http://localhost:8000/products

Params

Auth

Headers (8)

Body ●

Pre-req.

Tests

Settings

raw ▾

JSON ▾

```
1  {  
2      "name": "Pear",  
3      "price": 18,  
4      "quantity": 4  
5  }
```

Dashboard

local ^

Brokers

Topics

Consumers

Topics / products.created

Produce Message

⋮

Overview

Messages

Consumers

Settings

Statistics

Seek Type

Offset

Offset

Partitions

All items are selected.

Key Serde

String

Value Serde

String

Clear all

Submit

Oldest First

Search

+ Add Filters

DONE ⏱ 8 ms ⏲ 141 Bytes 📄 1 messages consumed

Offset	Partition	Timestamp	Key	Preview	Value	Preview
0	0	2/25/2026, 20:54:23	3		{"type": "product.created", "product_id": 3, "name": "Pear", "price": "18", "quantity": 4, "created_at": "2026-02-25T19:54:22.899267+00:00"}	

Key Value Headers



```
"type": "product.created",
"product_id": 3,
"name": "Pear",
"price": "18",
"quantity": 4,
"created_at": "2026-02-25T19:54:22.899267+00:00"
```

Timestamp 2/25/2026, 20:54:23
Timestamp type: CREATE_TIMEKey Serde String
Size: 1 BytesValue Serde String
Size: 140 Bytes

Listen for change With Products



compose.yml

```
consumer:
  build: .
  env_file:
    - .env
depends_on:
  - kafka
command: [ "uv", "run", "python", "-u", "app/consumer/print.py" ]
```

Docs: <https://kafka-python.readthedocs.io/en/master/>

THANKS !

Do you have any questions?
kristoffer.johansson@sti.se

CREDITS: This presentation template was created by Slidesgo, including icons by Flaticon, and infographics & images by Freepik.